

JMT/MTP Evaluation-I

Wamique Zia
2023CSM1020



Project Description

1. Adaptive Allocation:
Cache Replacement Policies at L2/L3 level.
 - Under the mentorship of Gautam Hazari @ AMD.
2. Exploiting parallelism to enhance Memory Performance.
 - Under the mentorship of Dr. T V Kalyan @ IIT Ropar.

Project 1 : Adaptive Allocation

Cache Replacement Policies at L2/L3 level.

- The basic idea relies on the 'Set Dueling' concept.
- Compares the performance of two policies on 'leader sets' and winning policy is applied to remaining sets called 'follower sets'.

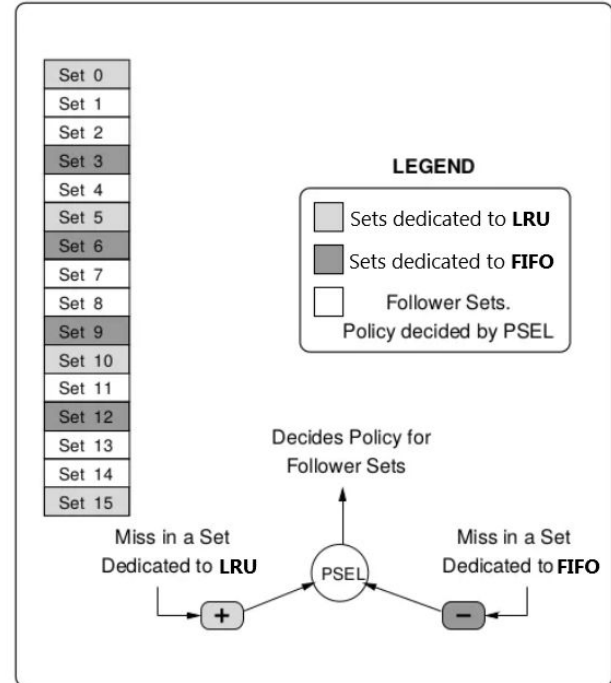


Fig. 1. Set Dueling b/w LRU and FIFO



Project 2 : **Exploiting parallelism to enhance Memory Performance**

1. Introduction
2. Literature Survey.
3. Trace collection for Graph and Genomics Applications.
4. DDR Frequency Sensitivity Analysis.



Project 2 : Introduction

- Why do we require it ?
 - ◆ Increased Memory Footprint
 - ◆ Cache Locality Limitations
 - ◆ Performance Improvement
 - ◆ Energy Efficiency

Project 2 : Literature Survey

- Different ideas/ways to enhance memory performance.
- ◆ Dynamic ROW ACTivation
 - ◆ DRAM Row Cache
 - ◆ Employing NoC packet routing mechanisms
 - ◆ Row/Page Management Policies
 - ◆ Perceptron Learning






Fig. 2. Conventional DRAM Organization

Dynamic/Partial ROW ACTivation


- Dynamic Detection of Row Size
 - ◆ Monitor Access Behavior, and
 - ◆ History Information Maintenance
- Calculating the Permutation Rate
 - ◆ Access Patterns
 - ◆ Granularity of Tracking
- Dynamic Activation Process
 - ◆ Row Activation
 - ◆ Early Precharge



(a) A 1KB memory row in terms of bytes.



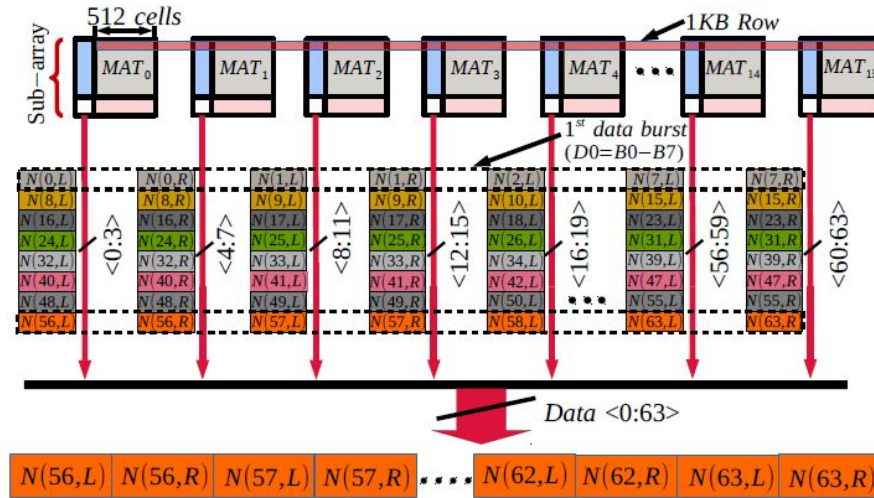
(b) A 1KB memory row in terms of columns.



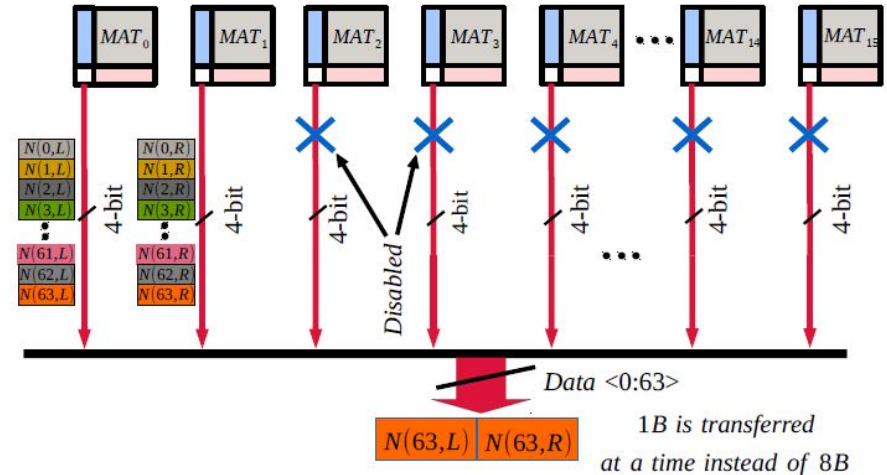
(c) A 1KB memory row divided into 4 equal-sized segments.

Fig. 3. 1KB memory row (page)

Dynamic/Partial ROW ACTivation (contd.)



(a) Conventional DRAM memory system.



(b) DRAM memory system that employs FGA_(1/8Bank) mechanism.

Fig. 4. An example shows the data fetching and accommodating in conventional DRAM memory systems and proposed DRA mechanism.



Dynamic/Partial ROW ACTivation (contd.)

→ Result:

- ◆ The Dynamic Row Activation (DRA) mechanism was tested using workloads from the MediaBench and PARSEC benchmark suites.
- ◆ DRA improves DRAM access latency by 12.3%.
- ◆ Achieves an average 39% reduction in ACTivation power.
- ◆ 12.8% reduction total DRAM energy consumption.

DRAM Row Cache

- The Row Buffer Cache (RBC) architecture is designed to cache the contents of the row buffer in an SRAM cache located within the memory controller. This allows the system to maintain access to frequently used rows, even when conflicts occur.

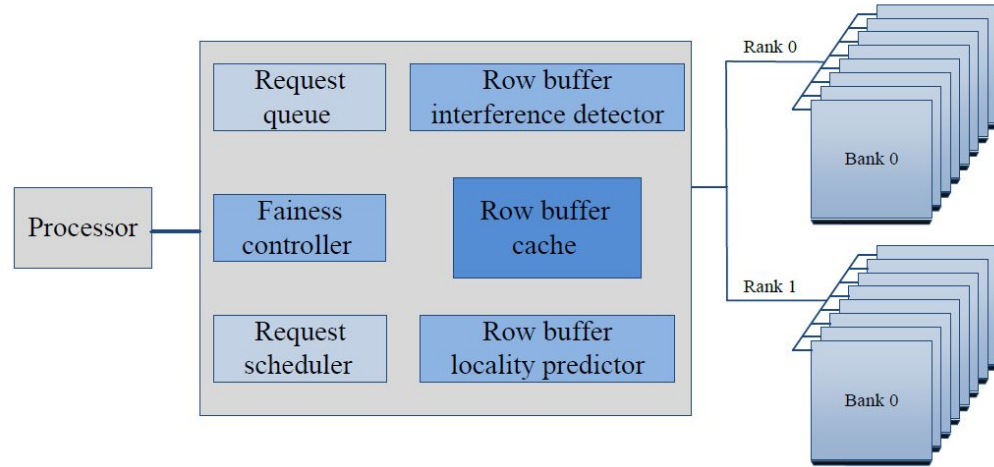


Fig. 5. Simplified View of Row Buffer Cache (RBC)



DRAM Row Cache (contd.)

- **Row Buffer Interference Detector:** This component quantifies the severity of row buffer interference for each row, identifying which rows are most affected by conflicts.
- **Row Buffer Locality Predictor:** This predictor uses the interference data to determine which rows have good locality and should be cached in the RBC.
- **Fairness Controller:** This ensures that all applications receive fair access to the RBC, preventing any single application from monopolizing the cache.

DRAM Row Cache (contd.)

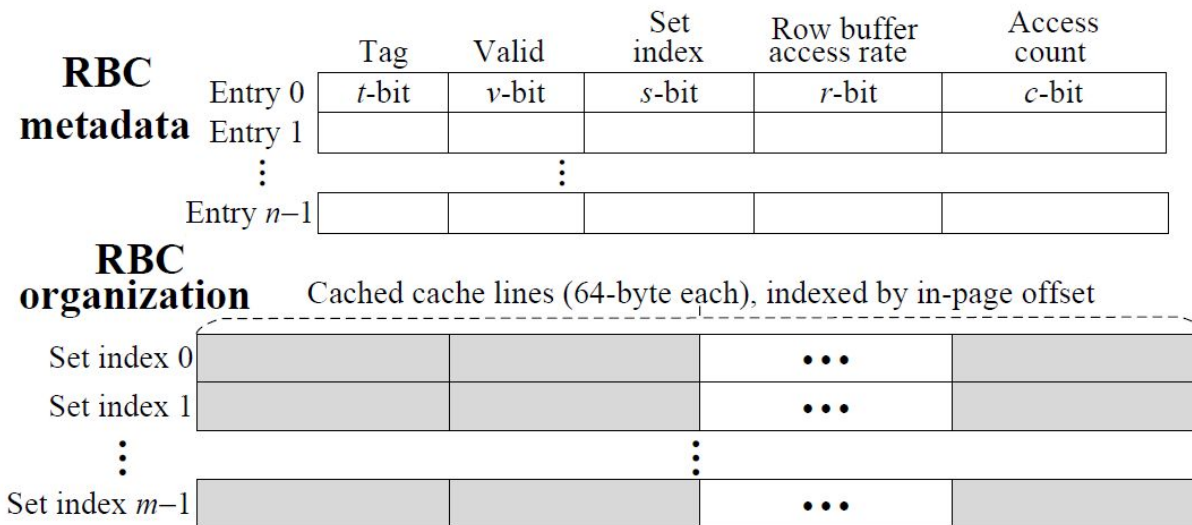


Fig. 6. Organization of RBC



DRAM Row Cache (contd.)

→ Result:

- ◆ In single-core simulations, RBC achieved up to a 2.24x performance improvement (16.1% on average).
- ◆ In multi-core simulations, performance increased by up to 1.55x (17% on average).
- ◆ Memory energy consumption was reduced by up to 68.2% (23.6% on average) in single-core simulations, and up to 35.4% (21.3% on average) in multi-core simulations.

Employing NoC packet routing mechanisms

- Challenges with heterogeneous MPSoCs:
- ◆ Memory Locality Dilution
 - ◆ Inadequate QoS Support
 - ◆ Increased Latency in Network-on-Chip (NoC)

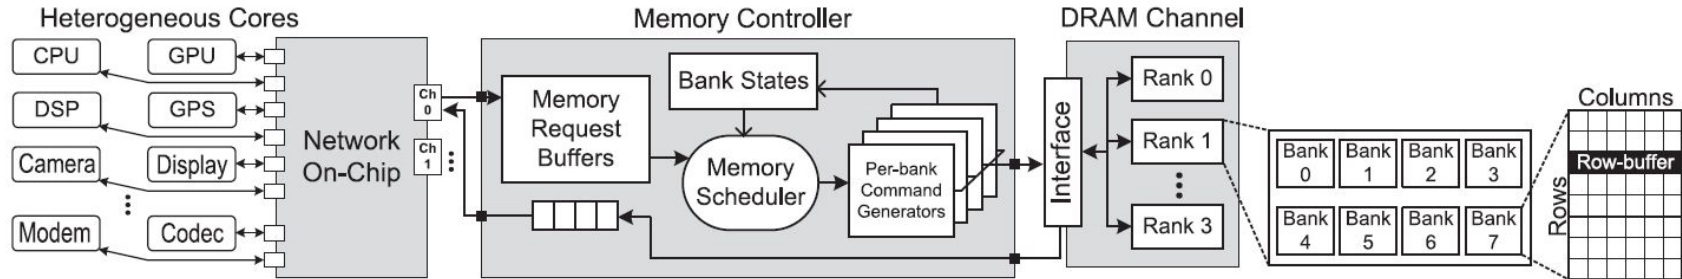


Fig. 7. Architecture of memory subsystem with a NoC for heterogeneous MPSoC system

Employing NoC packet routing mechanisms (contd.)

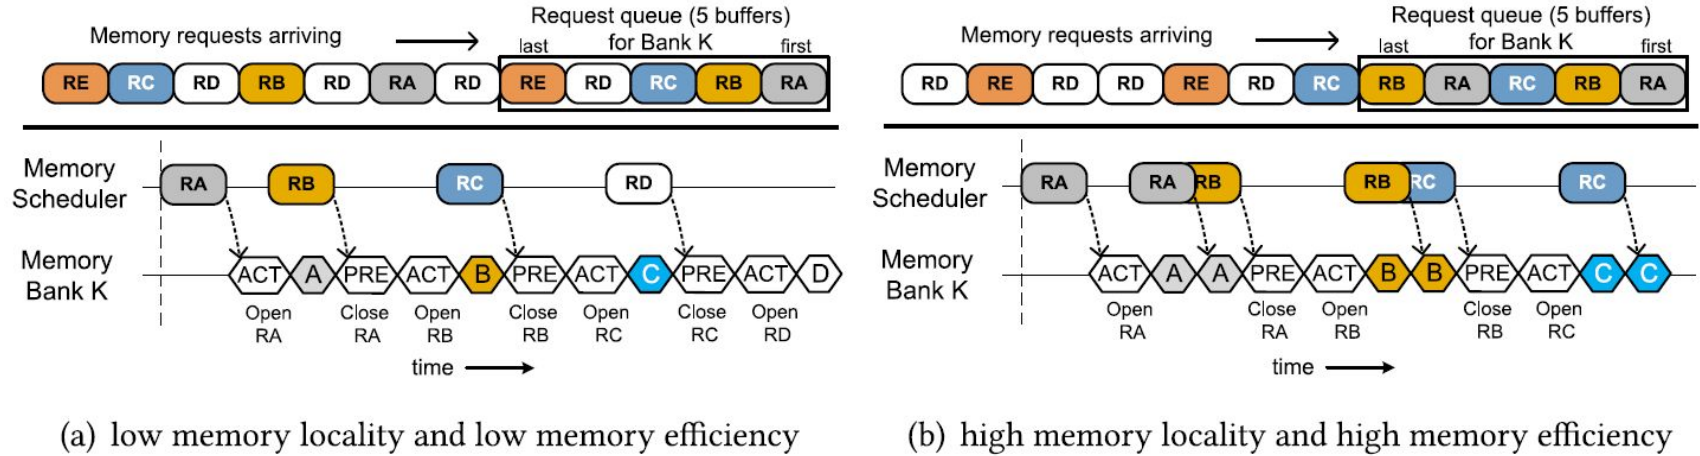


Fig. 8. An example of memory locality of NoC traffic affecting memory efficiency in DRAM



Employing NoC packet routing mechanisms (contd.)

→ Locality-aware Forwarding Mechanism

- ◆ Request Handling: When a memory request arrives at the NoC router, the router examines the row-index cache to determine if there are other requests for the same row.
- ◆ Forwarding Strategy: Router uses a locality-aware allocation scheme that considers both the row indices and the QoS requirements of the requests.
- ◆ Latency Reduction: Ensures that requests for the same row are forwarded in close succession



Employing NoC packet routing mechanisms (contd.)

→ Result:

- ◆ The locality-aware NoC router achieves a 58.32% increase in Row-Buffer Hits.
- ◆ The average memory latency is reduced by 14.45%.
- ◆ The design also results in a net reduction in energy costs for both DRAM and NoC operations by 27.82%.

Row/Page Management Policies

→ Problems:

- ◆ Fallback Allocation
- ◆ Promotion Failure
- ◆ Migration to CPU-less Nodes
- ◆ Inefficient Page Reclamation

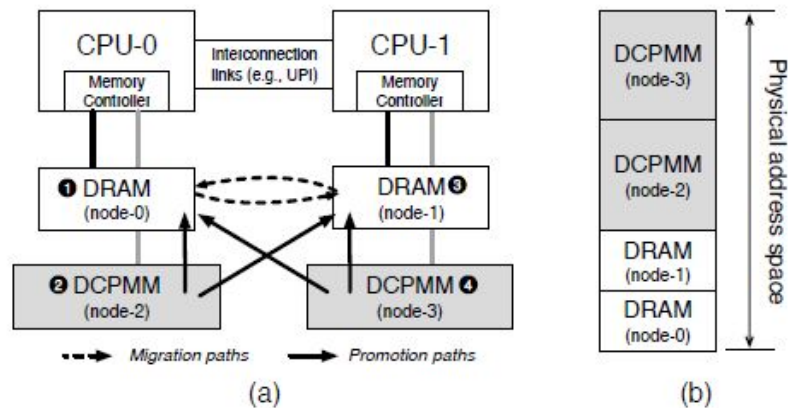


Fig. 9. Multi-tiered memory system augmented on the NUMA architecture



Row/Page Management Policies

- Auto-Tiering :
 - ◆ Promotion and Migration
 - CPM (Conservative Promotion and Migration)
 - OPM (Opportunistic Promotion and Migration)
 - ◆ Latency Hiding
 - ◆ Reclamation



Row/Page Management Policies

→ Results:

- ◆ Graph500 and GraphMat show performance increase of 17% and 19%, respectively.
- ◆ SPECAccel Workloads shows mixed results, with some sub-tests having 7-8% performance increase.
- ◆ OPM with Background Demotion provided up to a 6.17× speedup.
- ◆ OPM-BD outperformed Intel's Tiering approach by up to 6.99× in workloads like Graph500 and GraphMat.

Perceptron Learning

- Traditional page management policies (i.e., static open- or close-page policies) are unable to adapt to dynamically changing memory behaviors, leading to inefficiencies. Existing dynamic page policies are either access-based or time-based but still lack the ability to fully adapt to long-term memory access behavior, which changes over time.

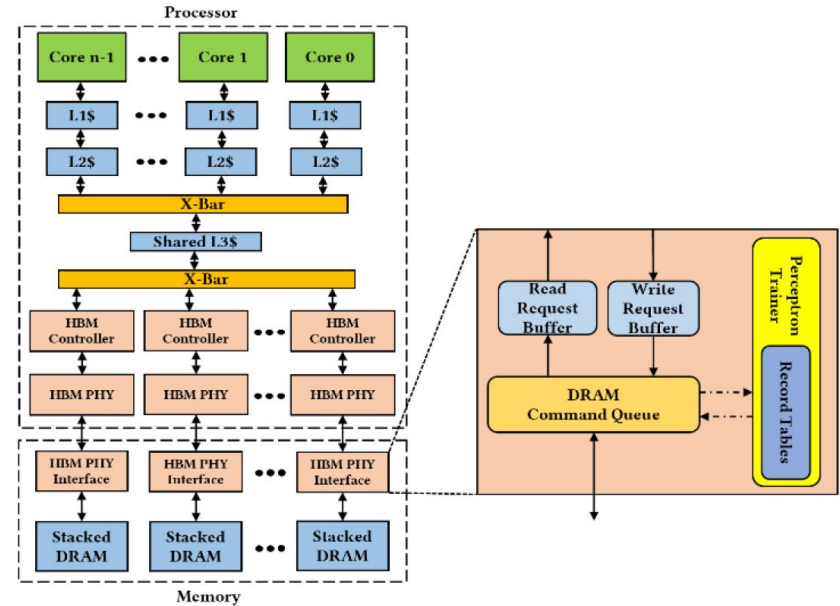


Fig. 10. Perceptron Trainer in HBM system



Perceptron Learning

- **DYMPL**, a perceptron-based dynamic page policy that adapts to the changing memory access behavior of applications by learning from past access patterns.
- Working:
 - ◆ Feature Extraction
 - ◆ Prediction
 - ◆ Training
 - ◆ Adaptation



Perceptron Learning

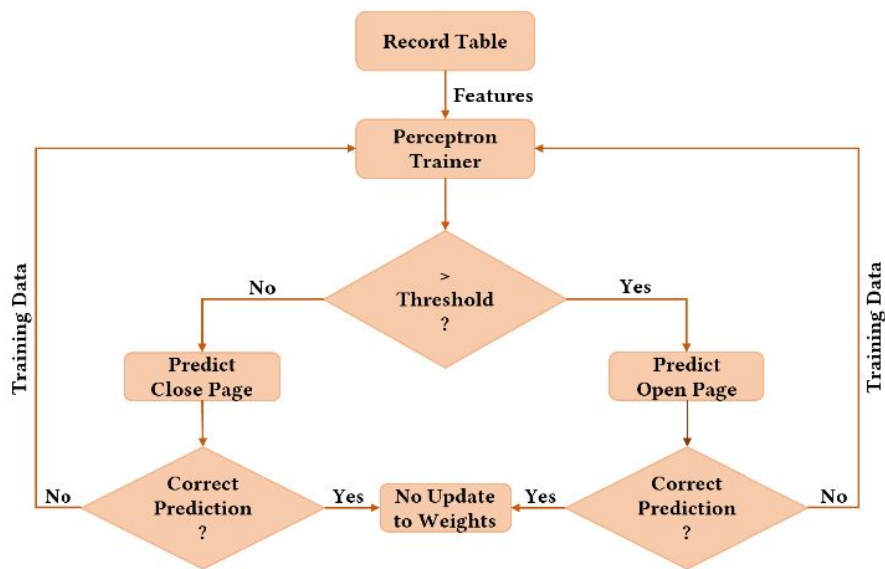


Fig. 11. Perceptron Training



Perceptron Learning

→ Results:

- ◆ DYMPL reduced row-buffer conflicts by 19.4% on average compared to static open-page policies.
- ◆ Overall system performance improved by 5.5% on average.
- ◆ Compared to statically implemented open- and close-page policies, the perceptron-based approach achieved performance improvements of 12.6% and 19.8%, respectively.



Trace Collection

1. GAPBS
 - BFS, PR, SSSP
2. GenomicsBench
 - BSW, CHAIN, DBG, PILEUP
3. Used Intel Pintool for tracing.

```
0x558898ebc7c3 R
0x7f69023f6f80 R
0x7f68eecf0db0 R
0x7f69023ffda0 W
0x7f68eeeb44ac R
0x7f68eeeba257 R
0x7f68eeeb3e94 R
0x7f68eeeb94c9 R
0x7f68eeeb3b7c R
0x7f68eeeb8d86 R
0x7ffcb0f05db8 W
0x7f68eeeb3cfc R
0x7f68eeeb90f1 R
0x7f68eeeb3dbc R
0x7f6902401d98 W
0x7f68eeeb92e3 R
```

Fig. 12. Sample Memory-Driven Trace File

Frequency Sensitivity Analysis




Chart. 1. DDR3 - Genomics/bsw

Frequency Sensitivity Analysis (contd.)

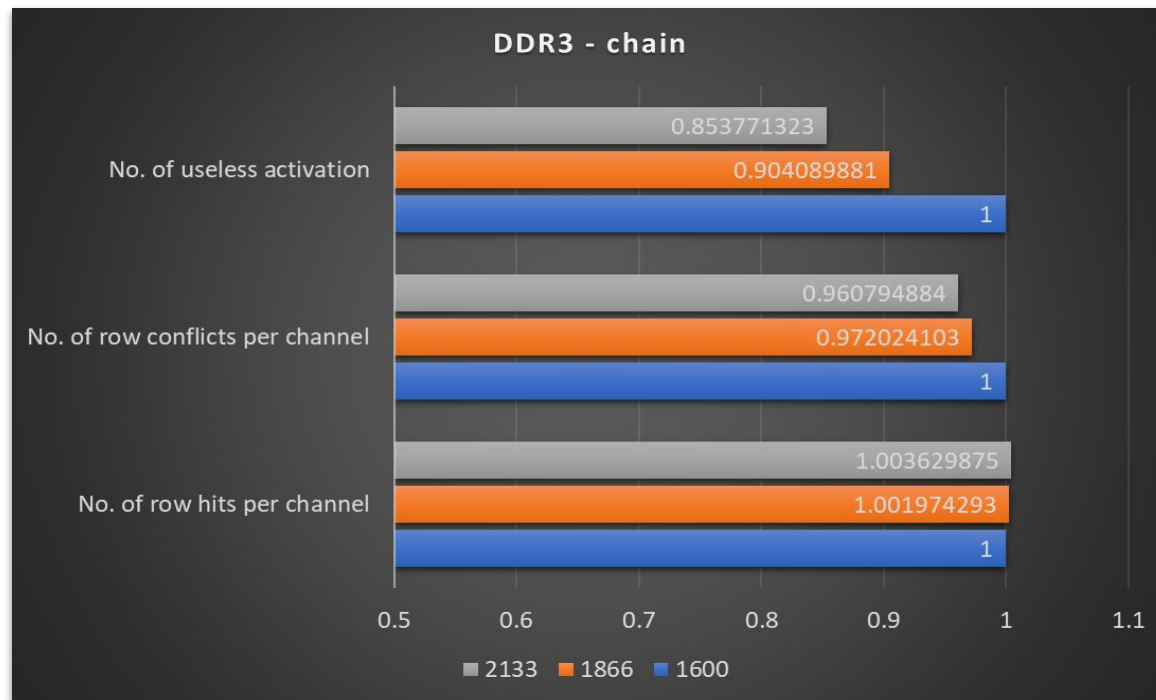


Chart. 2. DDR3 - Genomics/chain

Frequency Sensitivity Analysis (contd.)

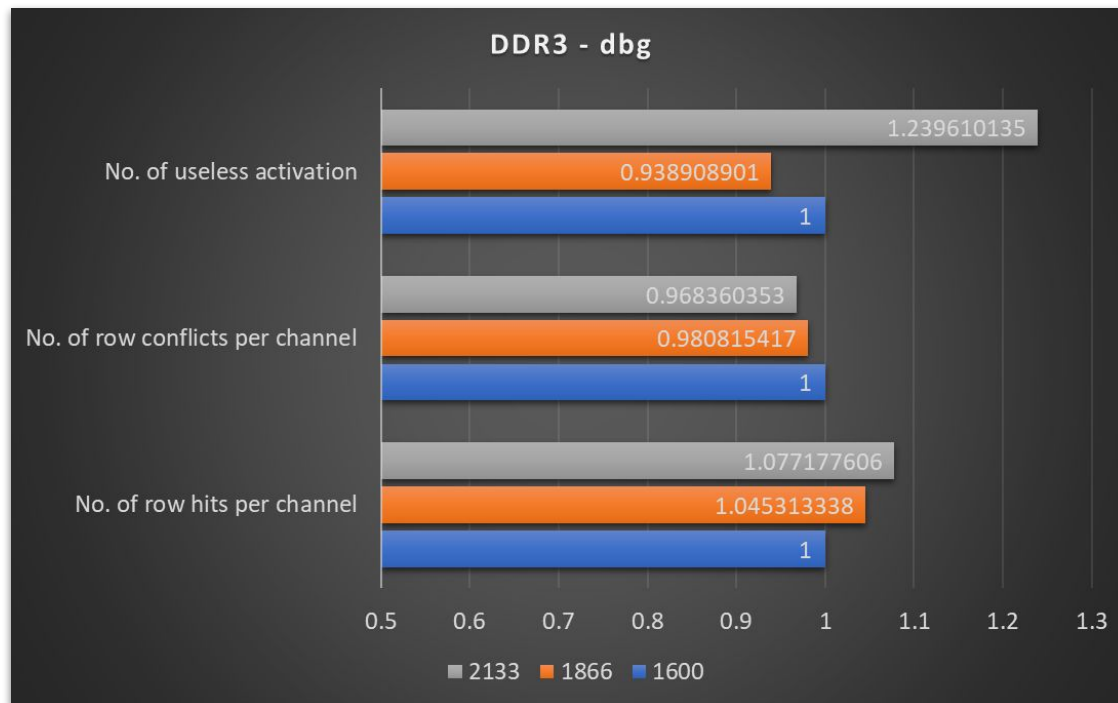


Chart. 3. DDR3 - Genomics/dbg

Frequency Sensitivity Analysis (contd.)




Chart. 4. DDR3 - GAP/bfs_amazon



Tools Used

1. **Ramulator** by CMU Safari group.
 - a. It is a fast and cycle-accurate DRAM simulator.
 - b. Supports memory-driven and cpu-driven trace.
 - c. Supports various DRAM standards, such as, DDR3, DDR4, LPDDR3, LPDDR4, GDDR5, and many more.



Future Work

1. Trace Collection for more applications.
2. Data bus width sensitivity analysis.
3. Per Column detailed study of memory access patterns.

Thank-You